



# Chapter : Threads



# Chapter : Threads

- Overview
- Benefits of Threads vs Processes
- Single and Multithreaded Processes
- Types of Threads
- Multithreading Models
- Thread Libraries

# Threads

- It is single sequential (flow of) execution of tasks of process
- A *thread* (or *lightweight process*) is a basic unit of CPU utilization; it consists of:
  - program counter
  - Thread id
  - register set
  - stack space
- A thread shares with its peer threads its:
  - code section
  - data section
  - operating-system resources


# Threads (Cont.)

- In a multiple threaded task, while one thread is blocked and waiting, a second thread in the same task can run.
  - In web browser, one thread displays images, other text, other fetches data from network.
  - Word processor, graphics, response to keystrokes, spell check.

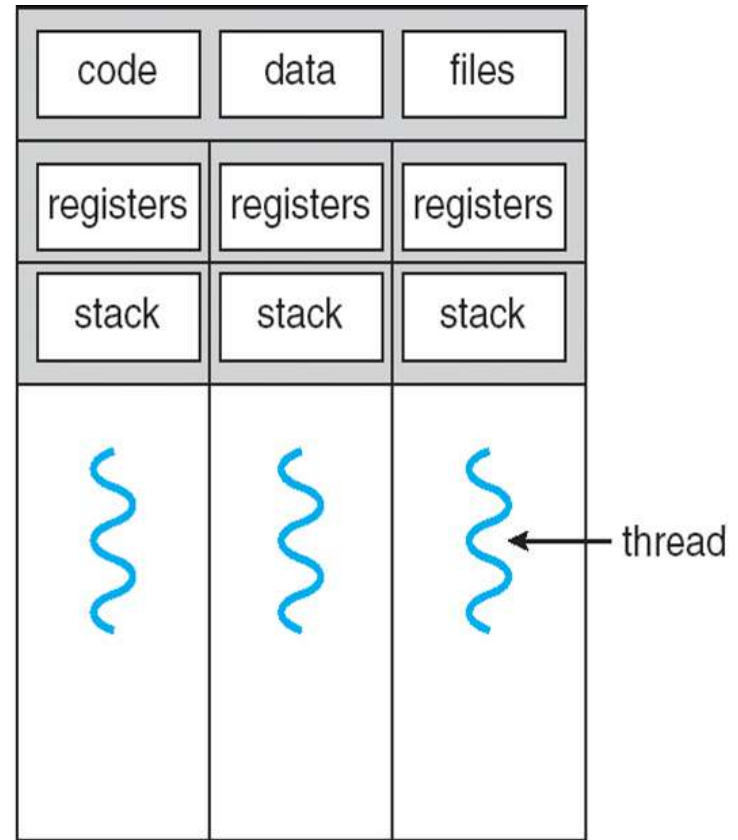
# Benefits of Threads vs Processes

- **Less time to create a new thread** than a process, because the newly created thread **uses the current process address space**.
- **Less time to terminate a thread**
- **Less time to switch between two threads** - newly created thread uses the current process address space.
- Less communication overheads - **threads share everything:** address space, So, **data produced by one thread is immediately available to all the other threads.**

# Single and Multithreaded Processes



single-threaded process



multithreaded process

# Threads States

- Three key states: running, ready, blocked
- Termination of a process, terminates all threads within the process

# Types of Threads

## 1. User Level Threads: (ULT)

- Threads of user application process.
- ULT are supported above the kernel and managed without kernel support
- These are implemented in user space in main memory. And managed by user level library.
- The kernel is not aware of the existence of threads
- User Level Library is used for – Thread Creation, Scheduling and Management



# Threads (Cont.)

- ULT require a kernel system call to operate
- It only takes care of the execution part.
- The lack of cooperation between user level threads and the kernel is a known disadvantage.
- In this case, the kernel may not favor a process that has many threads.
- User level threads are typically fast. Creating threads, switching between threads and synchronizing threads only needs a procedure call.
- If one thread blocks cause the entire process to block.

# User Level Threads

## **Advantages**

- Thread switching does not require Kernel mode privileges.
- User level thread can run on any operating system.
- Scheduling can be application specific in the user level thread.
- User level threads are fast to create and manage.

## **Disadvantages**

- In a typical operating system, most system calls are blocking.
- Multithreaded application cannot take advantage of multiprocessing.

# Kernel Level Threads(KLT)

- Threads of processes defined by OS itself
- KLT are supported and managed directly by OS.
- Kernel performs Thread Creation, Scheduling and Management in kernel space.
- No thread library but system calls to the kernel facility exists.



# Kernel Level Threads(KLT)

- Kernel level threads are managed by the OS
- thread operations (ex. Scheduling) are implemented in the kernel code.
- Kernel level threads may favor thread heavy processes.
- they can also utilize multiprocessor systems by splitting threads on different processors or cores.
- If one thread blocks it does not cause the entire process to block.
- KLT are not portable because the implementation is operating system dependent.

# Kernel Level Threads(KLT)

## **Advantages**

- Kernel can simultaneously schedule multiple threads from the same process on multiple processes.
- If one thread in a process is blocked, the Kernel can schedule another thread of the same process.
- Kernel routines themselves can be multithreaded.

## **Disadvantages**

- Kernel threads are generally slower to create and manage than the user threads.
- Transfer of control from one thread to another within the same process requires a mode switch to the Kernel.

# Combined ULT/KLT Approaches

- Thread creation done in the user space
- In a combined system, multiple threads within the same application can run in parallel on multiple processors and a blocking system call need not block the entire process
- The programmer may adjust the number of KLTs
- Example is Solaris

Lightweight processes (LWP) each LWP supports one or more ULTs and maps to exactly one KLT



# Multithreading Models

- One-to-One
- Many-to-One
- Many-to-Many

# One-to-One

- Each user-level thread maps to one kernel thread
- The one-to-one model associates a single user-level thread to a single kernel-level thread.
- kernel level threads follow the one to one model

Advantage:


facilitates the running of multiple threads in parallel.

Drawback:

Generation of every new user thread must include the creation of a corresponding kernel thread causing an overhead




# One-to-one Model



**The one-to-one model allows for greater concurrency, but the developer has to be careful not to create too many threads within an application**

# Many-to-One

- The many-to-one model associates all user-level threads to a single kernel-level thread.
- User level threads follow the many to one threading model.
- This means multiple threads managed by a library in user space but the kernel is only aware of a single thread of the process owning these threads.




# Many-to-Many Model

- In this model, developers can create as many user threads as necessary and the corresponding Kernel threads can run in parallel on a multiprocessor machine.
- This model provides the best accuracy on concurrency and when a thread performs a blocking system call, the kernel can schedule another thread for execution.

# Many-to-Many Model

- A number of user-level threads are associated to an equal or smaller number of kernel-level threads.
- Allows many user level threads to be mapped to many kernel threads

# Many-to-Many Model





# Thread Cancellation

- **Thread cancellation** is the task of terminating a thread before it has completed.
  - For example, if multiple threads are concurrently searching through a database and one thread returns the result, the remaining threads might be canceled.



# End of Chapter

# **Chapter : Process Synchronization**





# Process Synchronization

- Background
- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Semaphores
- Classic Problems of Synchronization
- Monitors
- Synchronization Examples
- Atomic Transactions

# Background

- **Co-Operating Process:** that can affect or be affected by other processes executing in system
- Concurrent access to shared data may result in data inconsistency
- **Process Synchronization:** Ensures coordination among processes and maintains Data Consistency

## ■ Process P1

1.  $X=5$
2.  $X=5+2$

3. `Printf( x);`

## Process P2

1. `read(x);`
2.  $x=x+5;$



# Producer Consumer Problem

There can be two situations:

- 1. Producer Produces Items at Fastest Rate Than Consumer Consumes**
- 2. Producer Produces Items at Lowest Rate Than Consumer Consumes**



# Producer Consumer Problem

1. **Producer Produces Items at Fastest Rate Than Consumer Consumes:**

If Producer produces items at fastest rate than Consumer consumes Then Some items will be lost


Eg. Computer → Producer  
Printer → Consumer

# Producer Consumer Problem

## Solution:

To avoid mismatch of items Produced or Consumed →  
Take Buffer

Idea is: Instead of sending items from Producer to  
Consumer directly → Store items into buffer





# Producer Consumer Problem

Buffer Can be:

## 1. Unbounderd Buffer:

1. No buffer size limit
2. Any no. of items can be stored
3. Producer can produce on any rate, there will always be space in buffer

## 2. Bounderd Buffer:

1. Limited buffer size



# Producer Consumer Problem

Bounded Buffer:

If rate of Production  $>$  rate of Consumption:

Some items will be unconsumed in buffer

If rate of Production  $<$  rate of Consumption:

At some time buffer will be empty



# Producer

```
while (true) {  
  
    /* produce an item and put in  
    nextProduced */  
  
    while (count == BUFFER_SIZE)  
        ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) %  
    BUFFER_SIZE;  
    count++;  
}
```



# Consumer

```
while (true) {  
    while (count == 0)           // buffer empty  
    {  
        ; // do nothing  
    }  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count- -;  
  
    /* consume the item in  
nextConsumed  
}
```

count=0  $\rightarrow$  buffer is empty

### **Producer**

count=count+1

1. R1=count

2. R1=R1+1

3. Count = R1

### **Consumer**

count=count-1

4. R2=count

5. R2=R2-1

6. Count = R2

Order of execution is not defined: May  
be

1  $\rightarrow$  2  $\rightarrow$  4  $\rightarrow$  5  $\rightarrow$  3  $\rightarrow$  6

# Race Condition

- When multiple processes access and manipulate the same data at the same time, they may enter into a race condition.
- **Race Condition: When output of the process is dependent on the sequence of other processes.**
- Race Condition occurs when processes share same data

## ■ Process P1

1. reads  $i=10$
2.  $i=i+1 = 11$

**PI got INTERRUPTED**

3. Stores  $i=11$  in memory

## Process P2

1. P2 reads  $i=11$  from memory
2.  $i=i+1 = 12$
3. Stores 12 in memory

# Critical Section Problem

- Section of code or set of operations, in which process may be changing shared variables, updating common data.
- A process is in the critical section if it executes code that manipulate shared data and resources.
- Each process should seek permission to enter its critical section → **Entry Section**
- **Exit Section**
- **Remainder section:** Contains remaining code

# Structure of a process

**Repeat**

**{**

**// Entry Section**

**Critical Section**

**// Exit Section**

**Remainder Section**

**} until false.**

# Solution to: Critical Section

## 1. Mutual Exclusion:

It states that **if one process is executing in its critical section**, then no other process can execute in its critical section.

## 2. Bounded Wait:

It states that if a process makes a request to enter its critical section and before that request is granted, **there is a limit on number of times other processes are allowed to enter that critical section**.

$P_i$  made request for Critical Section

$P_j$  is already executing in its Critical Section.....  $P_i$  has to wait

# Solution to: Critical Section

## 3. Progress:

It states that process cannot stop other process from entering their critical sections, if it is not executing in its CS.

Decision of entry into Critical Section is made by → Processes in **Entry Section**

**How?      →→→ Set Lock in Entry Section**

Once process leaves C.S, Lock is released (i.e in Exit Section)



# Solution to: Critical Section

## 3. Progress:

- N processes share same Critical Section
- Decision about which process can enter the Critical Section is based on Processes which are interested to enter the C.S  
(i.e. Processes of Entry Section)
- Processes that do not want to execute in CS should not take part in decision





# S/w Approaches/Solutions For C.S

- **Consider a Shared Variable**, that can take 2 values 0 and 1

If (shared variable == 0)

P0 can enter to Critical Section

If (shared variable == 1)

P1 can enter to Critical Section

# S/w Approaches/Solutions For C.S

■ Let **turn=0**

P0	P1
<pre>While (1) {     While(turn !=0);// if true     process will be stuck in the loop     CS     turn= 1; }</pre>	<pre>While (1) {     While(turn !=1);// if true     process will be stuck in the loop     CS     turn= 0; }</pre>



# S/w Approaches/Solutions For C.S

- **Mutual Exclusion satisfied**

- **Progress**

Now, Turn =1 i.e. process P1 can enter C.S.

**But**

Suppose P1 does not want to enter C.S.

**So.. turn =1 is never set**

Strict alteration is there.

**Progress is not satisfied**

## 2nd S/W Approach

Instead of one shared variable.

Assume a flag can be true or false.

- If p0 wants to enter into CS it will set its flag true and when it doesn't want to enter flag will be false.

## 2nd S/W Approach

P0	P1
<pre>While(1) {     Flag[0]=true;     while(Flag[1]);     CS     Flag[0]=false; }</pre>	<pre>While(1) {     Flag[1]=true;     while(Flag[0]);     CS     Flag[1]=false; }</pre>

## 2nd S/W Approach

- Follows mutual exclusion
- Doesn't follow progress
- System may go into deadlock state.

# Peterson's Solution



- Two process solution
- The two processes share two variables:
  - int **turn**;
  - Boolean **flag[2]**
- The variable **turn** indicates whose turn it is to enter the critical section.
- The **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i] = true** implies that process  $P_i$  is ready!

# Algorithm for Peterson's solution

**P0**

```
While(1)
{
    Flag[0]=true;
    turn=1;
    while(turn==1 &&
flag[1]==true);
    CS
    Flag[0]=false;
}
```

**P1**

```
While(1)
{
    Flag[1]=true;
    turn=0;
    while(turn==0 &&
flag[0]==true);
    CS
    Flag[1]=false;
}
```



- Mutual exclusion satisfied.
- Progress satisfied.
- Bounded wait satisfied.

# Synchronization Hardware



## Hardware Solution to C.S.

- Many systems provide hardware support for critical section code
- Uniprocessors – could disable interrupts
  - Currently running code would execute without preemption
- Disabling interrupt in multiprocessor environment can be time consuming , as message is passed to all the processors.
- Modern machines provide special atomic hardware instructions
  - ▶ Atomic = non-interruptable
  - Either test memory word and set value
  - Or swap contents of two memory words

## ■ Interrupt Disabling

- 1. In Single CPU system, processes do not execute in parallel**
- 2. Process leaves control of CPU when it is interrupted.**
- 3. Solution is:**
  - 1. To have each process disable all interrupts just after entering to the critical section.**
  - 2. Re-enable interrupts after leaving critical section**

## ■ Interrupt Disabling

**Repeat**

**Disable interrupts**

**C.S**

**Enable interrupts**

**Remainder section**

## ■ Hardware instructions

1. Machines provide inst. That can read, modify and store memory word
2. Common inst. are:

1. **Test and Set**

**This inst. Provides action of testing a variable and set its value**

**It executes atomically**

**Test and Set instructions operates on single Boolean variable.**

# TestAndndSet Instruction



## ■ Definition:

```
boolean TestAndSet (boolean *target)
{
    boolean rv = target;
    *target = TRUE;
    return rv;
}
```

# Solution using TestAndSet



- Shared Boolean variable lock., initialized to false.
- Solution:

Test and Set definition	
<pre>boolean TestAndSet (boolean *lock) {     boolean init = lock;     lock= TRUE;     return init; }</pre>	<pre>while (true) {     while ( TestAndSet (&amp;lock ));          // critical section          lock = FALSE;      // remainder section  }</pre>

# Swap Instruction



## ■ Definition:

```
void Swap (boolean *a, boolean *b)
{
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```



# Solution using Swap



- Shared Boolean variable lock initialized to FALSE; Each process has a local Boolean variable key.
- Solution:

Definition	
<pre>void Swap (boolean *a, boolean *b) {     boolean temp = *a;     *a = *b;     *b = temp; }</pre>	<pre>while (true) {     key = TRUE;     while ( key == TRUE)     {         Swap (&amp;lock, &amp;key );     }     // critical section      lock = FALSE;      // remainder section }</pre>

# Bounded waiting mutual exclusion with TestAndSet

```
■ do
{ waiting[i] = true;
key = true;
while (waiting[i] && key)
key = test_and_set(&lock);
waiting[i] = false;
/* critical section */
j = (i + 1) % n;
while ((j != i) && !waiting[j])
j = (j + 1) % n;
if (j == i)
lock = false;
else waiting[j] = false; /* remainder section */
}
while (true);
```

# Semaphore



- Semaphore S – integer variable
- Two standard operations modify S: `wait()` and `signal()`
  - Originally called `P()` and `V()`
- Less complicated
- Can only be accessed via two indivisible (atomic) operations
  - `wait (S) {`
    - `while S <= 0`
    - `; // no-op`
    - `S--;`
  - `}`
  - `signal (S) {`
    - `S++;`
  - `}`

- We can use semaphore to solve various synchronization problem.
- Let P1 and P2 are two process
- P1 wants to execute S1 statement. And P2 wants to execute S2 statement.
- S1 is to be executed before S2

S1;

Signal(synch);

Wait(synch);

S2

# Semaphore as General Synchronization Tool



- **Counting** semaphore – integer value can range over an unrestricted domain
- **Binary** semaphore – integer value can range only between 0 and 1; can be simpler to implement
  - Also known as **mutex locks**
- Can implement a counting semaphore **S** as a binary semaphore
- Provides mutual exclusion
  - Semaphore **S**; // initialized to 1
  - wait (**S**);  
    Critical Section  
    signal (**S**);

# Semaphore Implementation



- Must guarantee that no two processes can execute `wait ()` and `signal ()` on the same semaphore at the same time.
- Thus, implementation becomes the critical section problem where the wait and signal code are placed in the critical section.
  - Could now have busy waiting in critical section implementation
    - ▶ But implementation code is short

# Semaphore Implementation with no Busy waiting



- With each semaphore there is an associated waiting queue. Each entry in a waiting queue has two data items:
  - value (of type integer)
  - pointer to next record in the list
  
- Two operations:
  - **block** – place the process invoking the operation on the appropriate waiting queue.
  - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue.

# Semaphore Implementation with no Busy waiting (Cont.)



- Implementation of wait:

```
wait (S){  
    value--;  
    if (value < 0) {  
        add this process to waiting queue  
        block(); }  
}
```

- Implementation of signal:

```
Signal (S){  
    value++;  
    if (value <= 0) {  
        remove a process P from the waiting queue  
        wakeup(P); }  
}
```



# Deadlock and Starvation



- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let **S** and **Q** be two semaphores initialized to 1

$P_0$	$P_1$
wait (S);	wait (Q);
wait (Q);	wait (S);
.	.
.	.
.	.
signal (S);	signal (Q);
signal (Q);	signal (S);

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.

# Classical Problems of Synchronization



- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

# Bounded-Buffer Problem



- $N$  buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
- Semaphore **full** initialized to the value 0
- Semaphore **empty** initialized to the value  $N$ .

# Bounded Buffer Problem (Cont.)



- The structure of the producer process

```
while (true) {  
  
    // produce an item  
  
    wait (empty);  
    wait (mutex);  
  
    // add the item to the buffer  
  
    signal (mutex);  
    signal (full);  
}
```

# Bounded Buffer Problem (Cont.)



- The structure of the consumer process

```
while (true) {  
    wait (full);  
    wait (mutex);  
  
    // remove an item from  buffer  
  
    signal (mutex);  
    signal (empty);  
  
    // consume the removed item  
  
}
```

# Readers-Writers Problem



- A data set is shared among a number of concurrent processes
  - Readers – only read the data set; they do **not** perform any updates
  - Writers – can both read and write.
- Problem – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time.
- Shared Data
  - Data set
  - Semaphore **mutex** initialized to 1.
  - Semaphore **wrt** initialized to 1.
  - Integer **readcount** initialized to 0.

# Readers-Writers Problem (Cont.)



- The structure of a writer process

```
while (true) {  
    wait (wrt) ;  
  
    //  writing is performed  
  
    signal (wrt) ;  
}
```

# Readers-Writers Problem (Cont.)




- The structure of a reader process

```
while (true) {  
    wait (mutex) ;  
    readcount ++ ;  
    if (readcount == 1) wait (wrt) ;  
    signal (mutex)  
  
    // reading is performed  
  
    wait (mutex) ;  
    readcount - - ;  
    if (readcount == 0) signal (wrt) ;  
    signal (mutex) ;  
}
```



# Dining-Philosophers Problem



- Shared data
  - Bowl of rice (data set)
  - Semaphore **chopstick [5]** initialized to 1

# Dining-Philosophers Problem (Cont.)



- The structure of Philosopher  $i$ :

```
While (true) {  
    wait ( chopstick[i] );  
    wait ( chopStick[ (i + 1) % 5] );  
  
    // eat  
  
    signal ( chopstick[i] );  
    signal ( chopstick[ (i + 1) % 5] );  
  
    // think  
  
}
```

# Problems with Semaphores



- Incorrect use of semaphore operations:
  - signal (mutex) .... wait (mutex)
  - wait (mutex) ... wait (mutex)
  - Omitting of wait (mutex) or signal (mutex) (or both)

# Monitors




- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- Only one process may be active within the monitor at a time

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }
    ...

    procedure Pn (...) {.....}

    Initialization code ( ....) { ... }
    ...
}
}
```

# Schematic view of a Monitor

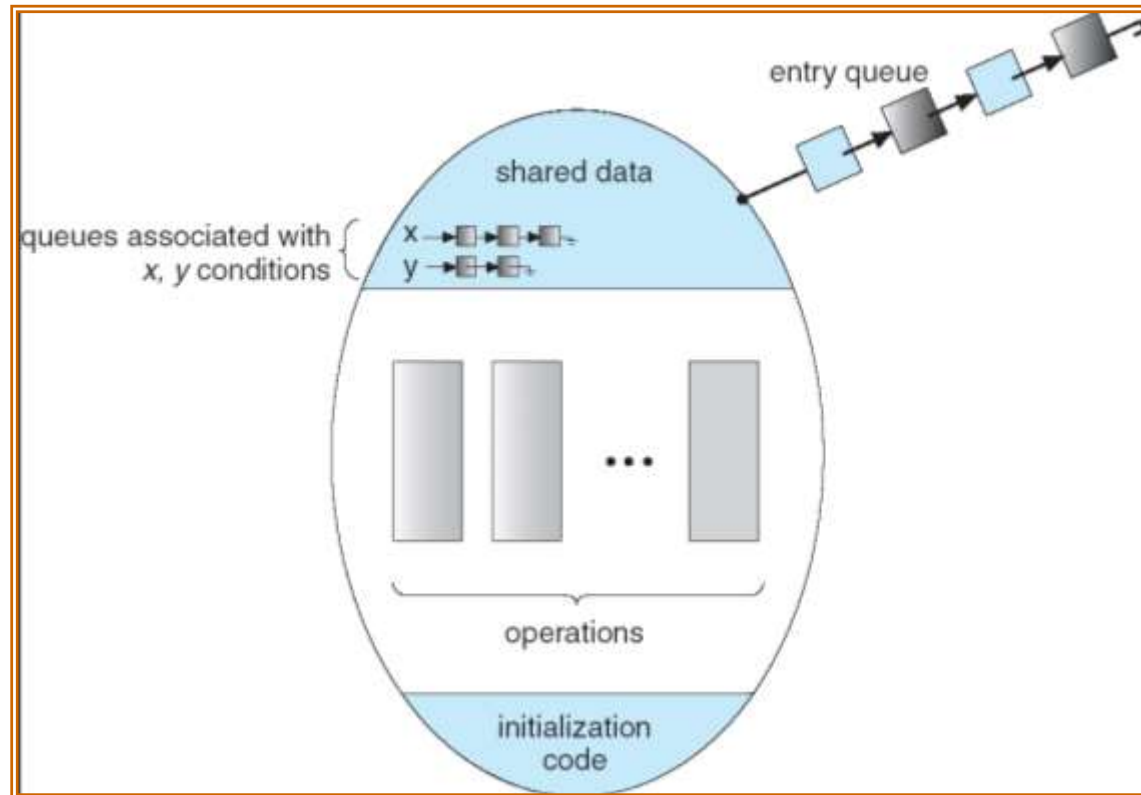


# Condition Variables



- condition x, y;
- Two operations on a condition variable:
  - x.wait () – a process that invokes the operation is suspended.
  - x.signal () – resumes one of processes (if any) that invoked x.wait ()

# Monitor with Condition Variables



# Solution to Dining Philosophers



monitor DP

```
{
    enum { THINKING; HUNGRY, EATING) state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING) self [i].wait;
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```



# Solution to Dining Philosophers (cont)



```
void test (int i) {  
    if ( (state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING ;  
        self[i].signal () ;  
    }  
}
```

```
    initialization_code() {  
        for (int i = 0; i < 5; i++)  
            state[i] = THINKING;  
    }  
}
```

# Solution to Dining Philosophers (cont)



- Each philosopher  $i$  invokes the operations `pickup()` and `putdown()` in the following sequence:

`dp.pickup (i)`

EAT

`dp.putdown (i)`

# Monitor Implementation Using Semaphores

- Variables

```
semaphore mutex; // (initially = 1)
semaphore next;  // (initially = 0)
int next-count = 0;
```

- Each procedure  $F$  will be replaced by

```
wait(mutex);
...
    body of  $F$ ;

...
if (next-count > 0)
    signal(next)
else
    signal(mutex);
```

- Mutual exclusion within a monitor is ensured.

# Monitor Implementation



- For each condition variable  $x$ , we have:

```
semaphore x-sem; // (initially = 0)
int x-count = 0;
```

- The operation  $x.\text{wait}$  can be implemented as:

```
x-count++;
if (next-count > 0)
    signal(next);
else
    signal(mutex);
wait(x-sem);
x-count--;
```



**End of Chapter 6**